

# Neural Networks: Part I


Daniel Yukimura

yukimura@impa.br

August 29, 2018

# Neural Networks I: Fundamentals


**Goal:** Learn a Parametric Function.



- $\theta \in \Theta$ : function parameters (**these are learned**).
- $\mathcal{X}$ : input space.
- $\mathcal{Y}$ : outcome space.

# The Perceptron


## *The Fundamental Building Block of Deep Learning*




# The Perceptron

## *The Fundamental Building Block of Deep Learning*

Processing units biologically inspired in **neurons**.



**Figure:** Neuron




**Figure:** (Artificial) Neuron

- There is no clear correspondence between Deep Learning and how the human brain works!

# The Perceptron

**Model:** A parametric function  $\phi : \mathbb{R}^k \rightarrow \mathbb{R}$ , given by




$$\phi(x) = \sigma \left( \sum_{i=1}^k w_i x_i + b \right)$$

- **activation function:**  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$  (usually non-linear).
- **parameters:**  $w = (w_1, \dots, w_k) \in \mathbb{R}^k$  and  $b \in \mathbb{R}$

# The Perceptron

**Model:** A parametric function  $\phi : \mathbb{R}^k \rightarrow \mathbb{R}$ , given by  
Is useful to look at it as a **feedforward flow!**




$$\phi(x) = \sigma \left( \sum_{i=1}^k w_i x_i + b \right)$$

- **activation function:**  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$  (usually non-linear).
- **parameters:**  $w = (w_1, \dots, w_k) \in \mathbb{R}^k$  and  $b \in \mathbb{R}$

# The Perceptron

**Model:** A parametric function  $\phi : \mathbb{R}^k \rightarrow \mathbb{R}$ , given by  
Is useful to look at it as a **feedforward flow!**




$$\phi(x) = \sigma \left( \sum_{i=1}^k w_i x_i + b \right)$$

- **activation function:**  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$  (usually non-linear).
- **parameters:**  $w = (w_1, \dots, w_k) \in \mathbb{R}^k$  and  $b \in \mathbb{R}$

# The Perceptron

**Model:** A parametric function  $\phi : \mathbb{R}^k \rightarrow \mathbb{R}$ , given by  
Is useful to look at it as a **feedforward flow!**




$$\phi(x) = \sigma \left( \sum_{i=1}^k w_i x_i + b \right)$$

- **activation function:**  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$  (usually non-linear).
- **parameters:**  $w = (w_1, \dots, w_k) \in \mathbb{R}^k$  and  $b \in \mathbb{R}$



# The Perceptron

**Model:** A parametric function  $\phi : \mathbb{R}^k \rightarrow \mathbb{R}$ , given by  
Is useful to look at it as a **feedforward flow!**




$$\phi(x) = \sigma \left( \sum_{i=1}^k w_i x_i + b \right)$$

- **activation function:**  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$  (usually non-linear).
- **parameters:**  $w = (w_1, \dots, w_k) \in \mathbb{R}^k$  and  $b \in \mathbb{R}$

# The Perceptron

**Model:** A parametric function  $\phi : \mathbb{R}^k \rightarrow \mathbb{R}$ , given by  
Is useful to look at it as a **feedforward flow!**




$$\phi(x) = \sigma \left( \sum_{i=1}^k w_i x_i + b \right)$$

- **activation function:**  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$  (usually non-linear).
- **parameters:**  $w = (w_1, \dots, w_k) \in \mathbb{R}^k$  and  $b \in \mathbb{R}$

# The Perceptron

## Reassessing Linear Models:


The addition of an **activation function** is the first step on rising model capacity.



# The Perceptron

## Reassessing Linear Models:


The addition of an **activation function** is the first step on rising model capacity.



# The Perceptron

## Reassessing Linear Models:

The addition of an **activation function** is the first step on rising model capacity.




# The Perceptron

## Common Activation Functions


### Rectified Linear Unit (ReLU)

$$\sigma(z) = \max(0, z)$$




### Sigmoid Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



### Hyperbolic Tangent



$$\sigma(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$



# The Perceptron

## Example: Binary Classification/Logistic Regression



The Perceptron was proposed as a model for **binary classification**. Originally it used the **step function** as activation.



# The Perceptron

## Example: Binary Classification/Logistic Regression

The Perceptron was proposed as a model for **binary classification**. Originally it used the **step function** as activation.



**Is hard to learn without differentiability!**




# The Perceptron

## Example: Binary Classification/Logistic Regression

In **logistic regression** we model the posterior distribution  $p(y | x)$  by smoothly squeezing the linear model into a probability distribution.

$$\begin{aligned} p_w(y = 1 | x) &= \text{sigm}(w^T x) \\ &= \frac{1}{1 + e^{-w^T x}} \end{aligned}$$



**meaning:** The probability that  $x$  belongs to the class 1.


# The Perceptron

## Example: The XOR function

- The Perceptron is unable to learn the exclusive or (XOR) function!
- The classes can't be separated by half-spaces (linear models).

$x_1$	$x_2$	$y$
0	0	0
0	1	1
1	0	1
1	1	0


Table:  $y = x_1 \oplus x_2$



# Neural Networks

## *How to combine neurons to build more expressive models?*

**Feedforward Neural Network (FNN):** We combine neurons layerwise as vertices of a directed graph.




- $h_j = \sigma \left( \sum_{i=1}^{n_0} w_{i,j}^{(1)} x_i + b_j \right)$
- $y_k = \sum_{j=1}^{n_1} w_{j,k}^{(2)} h_j$

# Neural Networks

## *How to combine neurons to build more expressive models?*

**Feedforward Neural Network (FNN):** We combine neurons layerwise as vertices of a directed graph.




- $h_j = \sigma \left( \sum_{i=1}^{n_0} w_{i,j}^{(1)} x_i + b_j \right)$
- $y_k = \sum_{j=1}^{n_1} w_{j,k}^{(2)} h_j$
- **Forward propagation**

# Neural Networks

## *How to combine neurons to build more expressive models?*

**Feedforward Neural Network (FNN):** We combine neurons layerwise as vertices of a directed graph.




- $h_j = \sigma \left( \sum_{i=1}^{n_0} w_{i,j}^{(1)} x_i + b_j \right)$
- $y_k = \sum_{j=1}^{n_1} w_{j,k}^{(2)} h_j$
- **Forward propagation**

# Neural Networks

## *How to combine neurons to build more expressive models?*

**Feedforward Neural Network (FNN):** We combine neurons layerwise as vertices of a directed graph.




- $h_j = \sigma \left( \sum_{i=1}^{n_0} w_{i,j}^{(1)} x_i + b_j \right)$
- $y_k = \sum_{j=1}^{n_1} w_{j,k}^{(2)} h_j$
- **Forward propagation**

# Neural Networks

## *How to combine neurons to build more expressive models?*

**Feedforward Neural Network (FNN):** We combine neurons layerwise as vertices of a directed graph.




- $h_j = \sigma \left( \sum_{i=1}^{n_0} w_{i,j}^{(1)} x_i + b_j \right)$
- $y_k = \sum_{j=1}^{n_1} w_{j,k}^{(2)} h_j$
- **Forward propagation**

# Neural Networks

## *How to combine neurons to build more expressive models?*

**Feedforward Neural Network (FNN):** We combine neurons layerwise as vertices of a directed graph.




- $h_j = \sigma \left( \sum_{i=1}^{n_0} w_{i,j}^{(1)} x_i + b_j \right)$
- $y_k = \sum_{j=1}^{n_1} w_{j,k}^{(2)} h_j$
- **Forward propagation**



# Neural Networks

## *How to combine neurons to build more expressive models?*

**Feedforward Neural Network (FNN):** We combine neurons layerwise as vertices of a directed graph.




- $h = \sigma \left( W^{(1)T} x + b \right)$
- $y = W^{(2)T} h$
- **Matrix notation is useful!**

# Neural Networks

## *How to combine neurons to build more expressive models?*


**Feedforward Neural Network (FNN):** We combine neurons layerwise as vertices of a directed graph.



- $h = \sigma \left( W^{(1)T} x + b \right)$
- $y = W^{(2)T} h$
- **Universal Approximation Theorem:** Given enough neurons in a hidden layer, and a non-linear increasing activation function, one can approximate any Borel measurable function (see [ref](#)).

# Neural Networks


## Do we need more layers?



- Using more layers seems to allow more capacity while using fewer neurons, see [\[ref\]](#).
- There are many cases of success by using more layers.
- Deeper networks are harder to train!

# Neural Networks


## Do we need more layers?



- Using more layers seems to allow more capacity while using fewer neurons, see [\[ref\]](#).
- There are many cases of success by using more layers.
- Deeper networks are harder to train!

# Neural Networks


## Do we need more layers?



- Using more layers seems to allow more capacity while using fewer neurons, see [ref].
- There are many cases of success by using more layers.
- Deeper networks are harder to train!

# Neural Networks

## Do we need more layers?



- $h^{(0)} = x, h^{(\ell)} = \sigma \left( W^{(\ell)T} h^{(\ell-1)} + b^{(\ell)} \right), \ell \in [L-1]$
- $\hat{y} = f(x, \theta) = W^{(L)T} h^{(L-1)}$  (sometimes  $\hat{y} = \sigma(\dots)$ ).
- We denote  $\theta_\ell = (W^{(\ell)}, b^{(\ell)})$  the parameters of layer  $\ell$ , and  $\theta = (\theta_1, \dots, \theta_L)$

# Risk Minimization

## Recall:

We want to find the network weights that **achieve the lowest risk value**.

$$\begin{aligned}\hat{\theta} &= \operatorname{argmin}_{\theta \in \Theta} R_n(\theta) \\ &= \operatorname{argmin}_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n \ell(f(x_i, \theta), y_i)\end{aligned}$$

**Example:** For  $L_2$  regression we have

$$R_n(\theta) = \frac{1}{n} \sum_{i=1}^n \|f(x_i, \theta) - y_i\|_2^2$$

# Maximum Likelihood Estimation

- When modelling posterior distributions  $p_{\theta}(y|x)$  is useful to look at the likelihood function

$$\mathcal{L}_n(\theta) = p_{\theta}(\mathcal{D}) = \prod_{i=1}^n p_{\theta}(x_i, y_i)$$

- Maximizing  $\mathcal{L}_n(\theta)$  means finding  $p_{\theta}$  that best represents the data.
- But, in the supervised problem we can consider the alternative form

$$\mathcal{L}_n(\theta) = \prod_{i=1}^n p_{\theta}(y_i | x_i) .$$



# Maximum Likelihood Estimation

- When modelling posterior distributions  $p_{\theta}(y|x)$  is useful to look at the likelihood function

$$\mathcal{L}_n(\theta) = p_{\theta}(\mathcal{D}) = \prod_{i=1}^n p_{\theta}(x_i, y_i)$$

- Maximizing  $\mathcal{L}_n(\theta)$  means finding  $p_{\theta}$  that best represents the data.
- But, in the supervised problem we can consider the alternative form

$$\mathcal{L}_n(\theta) = \prod_{i=1}^n p_{\theta}(y_i | x_i) .$$

# Maximum Likelihood Estimation

- When modelling posterior distributions  $p_{\theta}(y|x)$  is useful to look at the likelihood function

$$\mathcal{L}_n(\theta) = p_{\theta}(\mathcal{D}) = \prod_{i=1}^n p_{\theta}(x_i, y_i)$$

- Maximizing  $\mathcal{L}_n(\theta)$  means finding  $p_{\theta}$  that best represents the data.
- But, in the supervised problem we can consider the alternative form

$$\mathcal{L}_n(\theta) = \prod_{i=1}^n p_{\theta}(y_i | x_i) .$$

# Maximum Likelihood Estimation

- The negative log-likelihood translates into the risk problem

$$-\frac{1}{n} \log (\mathcal{L}_n(\theta)) = \frac{1}{n} \sum_{i=1}^n -\log p_{\theta}(y_i | x_i)$$

- Therefore, the Maximum Likelihood Estimator (MLE) can be obtained through minimizing such risk

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n -\log p_{\theta}(y_i | x_i)$$

# Maximum Likelihood Estimation

- The negative log-likelihood translates into the risk problem

$$-\frac{1}{n} \log (\mathcal{L}_n(\theta)) = \frac{1}{n} \sum_{i=1}^n -\log p_{\theta}(y_i | x_i)$$

- Therefore, the Maximum Likelihood Estimator (MLE) can be obtained through minimizing such risk

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n -\log p_{\theta}(y_i | x_i)$$

# Maximum Likelihood Estimation

## Example: Binary Classification/Logistic Regression

- Observe that in the binary classification case  $y_i \in \{0, 1\}$  we can write the posterior as

$$p_{\theta}(y_i | x_i) = f(x_i, \theta)^{y_i} (1 - f(x_i, \theta))^{(1-y_i)}$$

- Implying

$$\log p_{\theta}(y_i | x_i) = y_i \log(f(x_i, \theta)) + (1 - y_i) \log(1 - f(x_i, \theta))$$

# Maximum Likelihood Estimation

## Example: Binary Classification/Logistic Regression

- Observe that in the binary classification case  $y_i \in \{0, 1\}$  we can write the posterior as

$$p_{\theta}(y_i | x_i) = f(x_i, \theta)^{y_i} (1 - f(x_i, \theta))^{(1-y_i)}$$

- Implying

$$\log p_{\theta}(y_i | x_i) = y_i \log(f(x_i, \theta)) + (1 - y_i) \log(1 - f(x_i, \theta))$$

# Maximum Likelihood Estimation

## Example: Binary Classification/Logistic Regression

- The resulting loss is called the **Cross Entropy Loss**

$$R_n(\theta) = \frac{1}{n} \sum_{i=1}^n y_i \log(f(x_i, \theta)) + (1 - y_i) \log(1 - f(x_i, \theta))$$

- The next question is how to actually optimize such functions.

# Maximum Likelihood Estimation

## Example: Binary Classification/Logistic Regression

- The resulting loss is called the **Cross Entropy Loss**

$$R_n(\theta) = \frac{1}{n} \sum_{i=1}^n y_i \log(f(x_i, \theta)) + (1 - y_i) \log(1 - f(x_i, \theta))$$

- The next question is how to actually optimize such functions.




# Gradient Descent

The classical **gradient descent** (GD) consists on the iteration

$$\theta_{t+1} = \theta_t - \alpha \nabla L(\theta_t)$$

for some initial configuration  $\theta_0$  and learning rate  $\alpha > 0$ .



# Computing Gradients: Backpropagation

**Backpropagation** is an efficient algorithm for computing risk gradients of NN models (Is essentially just chain rule).

- Let  $L(\theta) = c(f(x, \theta))$ , where the cost function  $c$  might depend on the label  $y$  or other parameters, but for the derivation purpose they are omitted.
- How does a small change in the parameters  $\theta_\ell$  affect the loss  $L$ ?
- Observe that  $L(\theta) = L(h^{(\ell)}(h^{(\ell-1)}, \theta_\ell), \theta_{\ell+1}^L)$ , then

$$\frac{\partial L}{\partial \theta_\ell} = \sum_{j=1}^{|H_\ell|} \frac{\partial L}{\partial h_j^{(\ell)}} \cdot \frac{\partial h_j^{(\ell)}}{\partial \theta_\ell}$$

# Computing Gradients: Backpropagation

**Backpropagation** is an efficient algorithm for computing risk gradients of NN models (Is essentially just chain rule).

- Let  $L(\theta) = c(f(x, \theta))$ , where the cost function  $c$  might depend on the label  $y$  or other parameters, but for the derivation purpose they are omitted.
- How does a small change in the parameters  $\theta_\ell$  affect the loss  $L$ ?
- Observe that  $L(\theta) = L(h^{(\ell)}(h^{(\ell-1)}, \theta_\ell), \theta_{\ell+1}^L)$ , then

$$\frac{\partial L}{\partial \theta_\ell} = \sum_{j=1}^{|H_\ell|} \frac{\partial L}{\partial h_j^{(\ell)}} \cdot \frac{\partial h_j^{(\ell)}}{\partial \theta_\ell}$$

# Computing Gradients: Backpropagation

**Backpropagation** is an efficient algorithm for computing risk gradients of NN models (Is essentially just chain rule).

- Let  $L(\theta) = c(f(x, \theta))$ , where the cost function  $c$  might depend on the label  $y$  or other parameters, but for the derivation purpose they are omitted.
- How does a small change in the parameters  $\theta_\ell$  affect the loss  $L$ ?
- Observe that  $L(\theta) = L(h^{(\ell)}(h^{(\ell-1)}, \theta_\ell), \theta_{\ell+1}^L)$ , then

$$\frac{\partial L}{\partial \theta_\ell} = \sum_{j=1}^{|H_\ell|} \frac{\partial L}{\partial h_j^{(\ell)}} \cdot \frac{\partial h_j^{(\ell)}}{\partial \theta_\ell}$$

# Computing Gradients: Backpropagation

**Backpropagation** is an efficient algorithm for computing risk gradients of NN models (Is essentially just chain rule).

- Let  $L(\theta) = c(f(x, \theta))$ , where the cost function  $c$  might depend on the label  $y$  or other parameters, but for the derivation purpose they are omitted.
- How does a small change in the parameters  $\theta_\ell$  affect the loss  $L$ ?
- Observe that  $L(\theta) = L(h^{(\ell)}(h^{(\ell-1)}, \theta_\ell), \theta_{\ell+1}^L)$ , then

$$\frac{\partial L}{\partial \theta_\ell}_{\text{vector}} = \sum_{j=1}^{|H_\ell|} \frac{\partial L}{\partial h_j^{(\ell)}}_{\text{vector}} \cdot \frac{\partial h_j^{(\ell)}}{\partial \theta_\ell}_{\text{matrix}}$$

# Computing Gradients: Backpropagation

- Observe that  $L(\theta) = L(h^{(\ell)}(h^{(\ell-1)}, \theta_\ell), \theta_{\ell+1}^L)$ , then

$$\frac{\partial L}{\partial \theta_\ell}_{\text{vector}} = \sum_{j=1}^{|H_\ell|} \frac{\partial L}{\partial h_j^{(\ell)}}_{\text{vector}} \cdot \frac{\partial h_j^{(\ell)}}{\partial \theta_\ell}_{\text{matrix}}$$

- $\frac{\partial h_j^{(\ell)}}{\partial \theta_\ell}$  can be computed directly from the definition.

# Computing Gradients: Backpropagation

- Observe that  $L(\theta) = L(h^{(\ell)}(h^{(\ell-1)}, \theta_\ell), \theta_{\ell+1}^L)$ , then

$$\frac{\partial L}{\partial \theta_\ell}_{\text{vector}} = \sum_{j=1}^{|H_\ell|} \frac{\partial L}{\partial h_j^{(\ell)}}_{\text{vector}} \cdot \frac{\partial h_j^{(\ell)}}{\partial \theta_\ell}_{\text{matrix}}$$

- $\frac{\partial h_j^{(\ell)}}{\partial \theta_\ell}$  can be computed directly from the definition.

# Computing Gradients: Backpropagation

- The vector  $\delta_\ell = \frac{\partial L}{\partial h_j^{(\ell)}}$  can be computed through a recursion on the network, on the opposite direction, starting from  $L$ 
  - $\delta_L = \frac{\partial L}{\partial h^{(L)}}$  is just the gradient of the cost function  $c(\cdot)$ .
  - For  $\ell \in [L - 1]$

$$\delta_\ell = \frac{\partial L}{\partial h^{(\ell+1)}} \frac{\partial h^{(\ell+1)}}{\partial h^{(\ell)}} = \delta_{\ell+1} \frac{\partial h^{(\ell+1)}}{\partial h^{(\ell)}}$$

- The values of  $\frac{\partial h^{(\ell+1)}}{\partial h^{(\ell)}}$  can also be computed directly.



# Stochastic Gradient Descent (SGD)

- On each iteration  $t > 0$  we choose uniformly at random an  $S$ -set  $\mathcal{S} \subseteq [N]$  of indices ( $|\mathcal{D}| = N$ ) and compute the *minibatch* gradient as

$$\hat{L}_{\mathcal{S}}(\theta) = \frac{1}{S} \sum_{i \in \mathcal{S}} \ell(\theta, \mathbf{z}_i)$$

$$\hat{g}_{\mathcal{S}}(\theta) = \nabla \hat{L}_{\mathcal{S}}(\theta)$$

- The iteration is given as before

$$\theta_{t+1} = \theta_t - \alpha \hat{g}_{\mathcal{S}}(\theta)$$

# Stochastic Gradient Descent (SGD)

**Remark:** The noise resulting from working with minibatches actually helps on avoiding bad minimas and to escape saddle points.

